

DEEP LEARNING TOOLS and FRAMEWORKS

HAMID PALANGI

DEEP LEARNING GROUP, MICROSOFT RESEARCH AI

REDMOND, WA, USA

November 16, 2017 at IEEE GlobalSIP, Montreal, Canada

Acknowledgements: Xiaodong He, Susan Dumais, Li Deng, Jianfeng Gao, Kenneth Tran, Yelong Shen, Xinying Song, Po-Sen Huang, Paul Smolensky, Z. Jane Wang, Rabab Ward



DEEP LEARNING (DL)

- Is it “always” good to use DL models for my task?

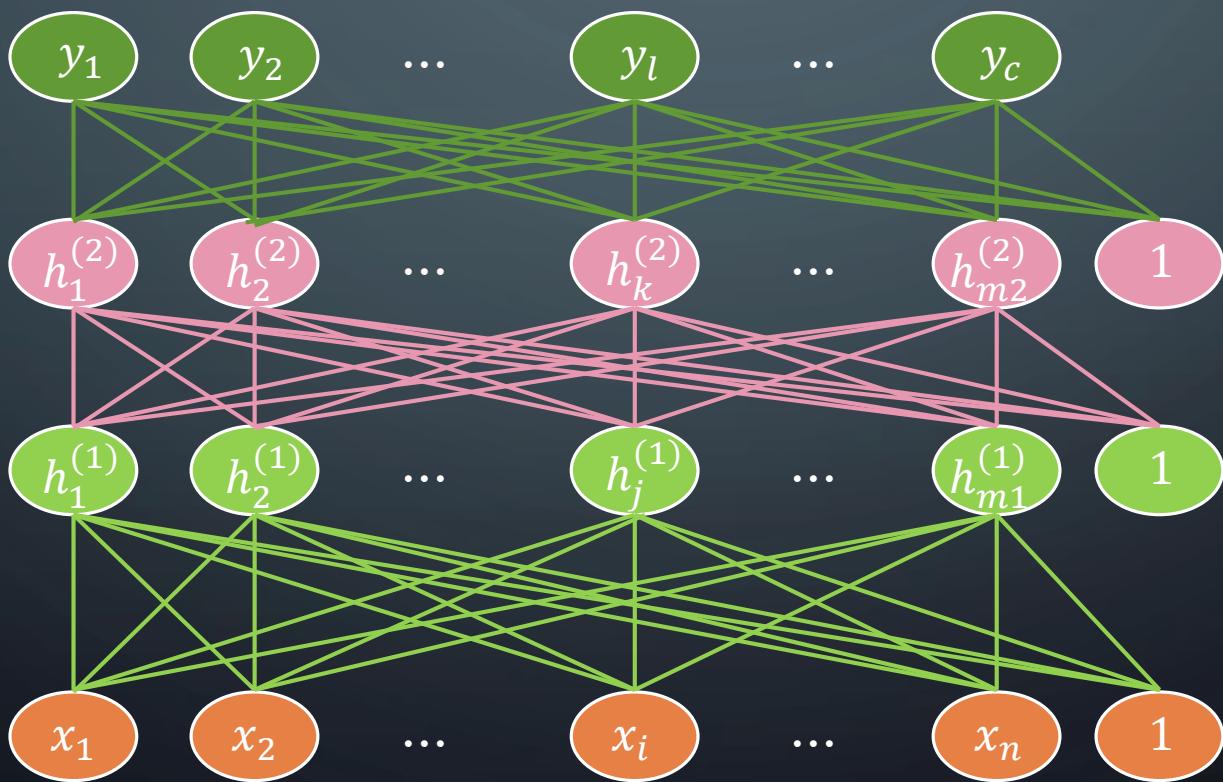
DEEP LEARNING (DL)

- Is it “always” good to use DL models for my task?

No!!!

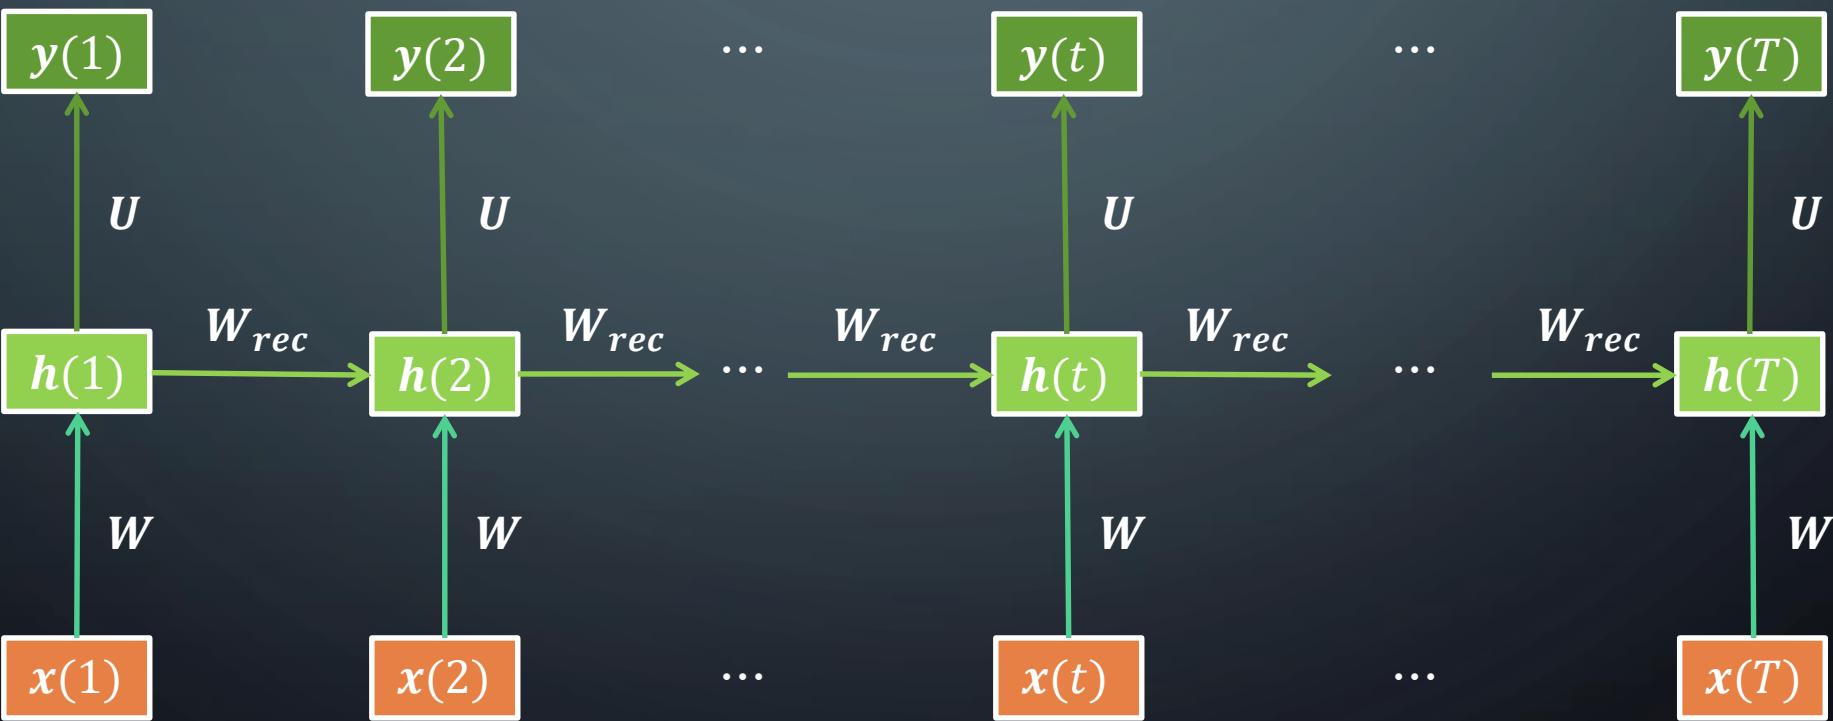
DEEP LEARNING (DL): SOME EXAMPLES

- Feedforward NN



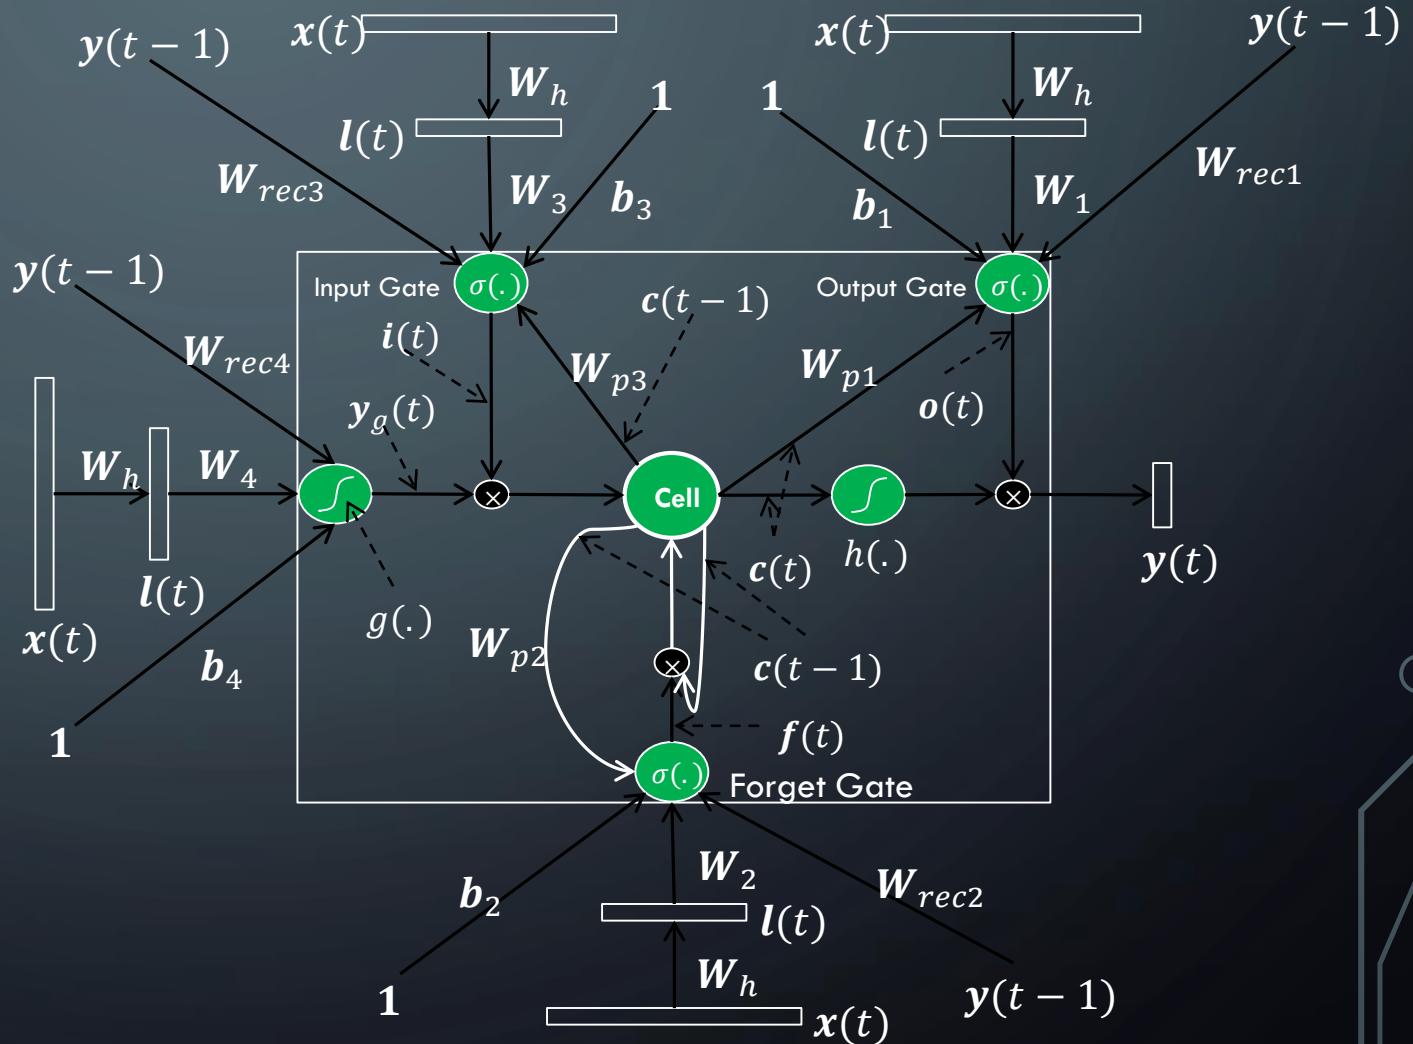
DEEP LEARNING (DL): SOME EXAMPLES

- Recurrent Neural Network (RNN)



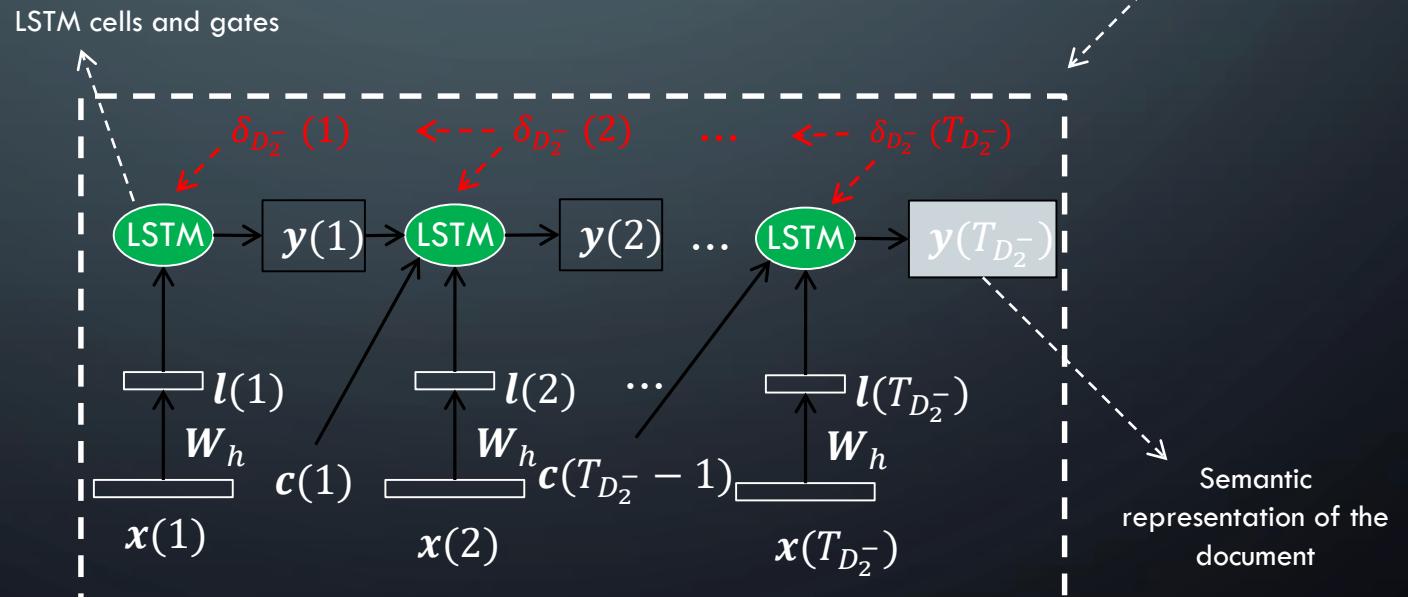
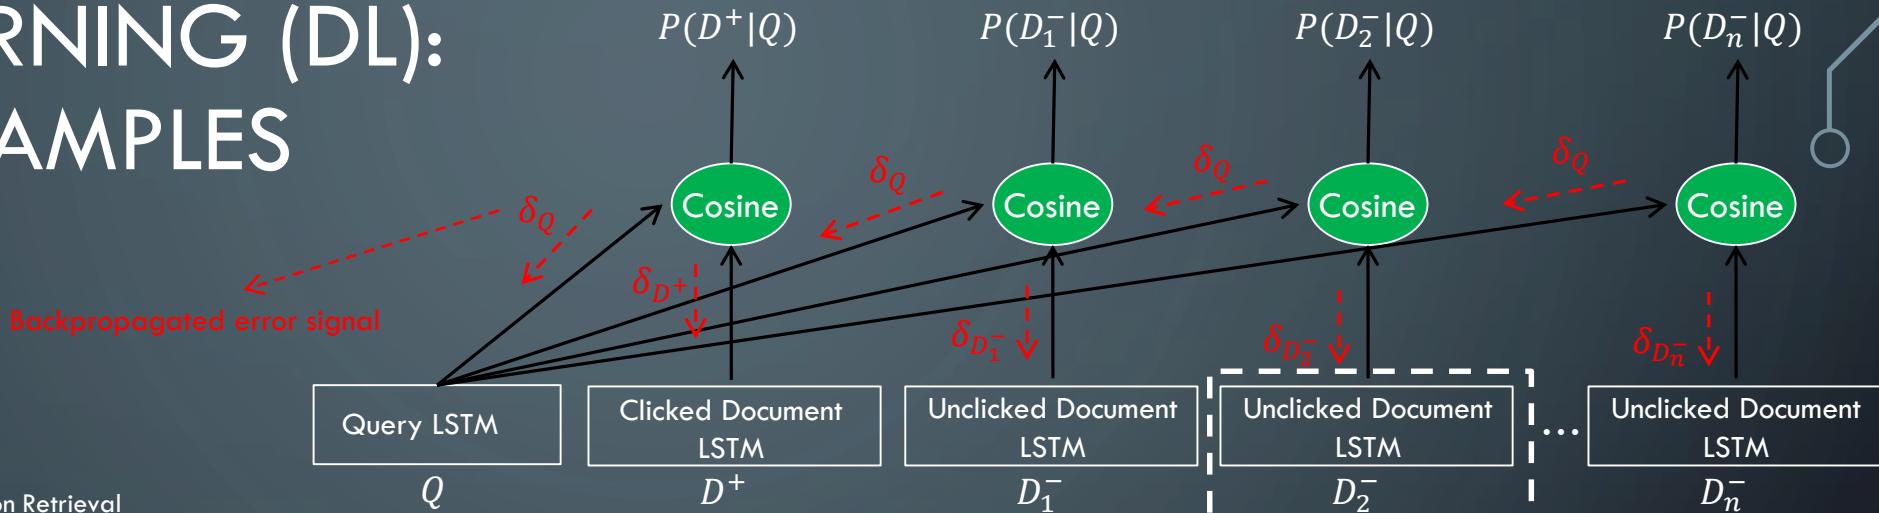
DEEP LEARNING (DL): SOME EXAMPLES

- Long Short-Term Memory
 - LSTM



DEEP LEARNING (DL): SOME EXAMPLES

LSTM-DSSM used in Information Retrieval
<https://arxiv.org/abs/1502.06922>



MANUAL GRADIENT CALCULATION

- Good idea to learn stuff. Bad idea to get the job done ASAP.

$$\frac{\partial \hat{y}(t)}{\partial w_{rec_3}} = \frac{\partial}{\partial w_{rec_3}} (\text{diag}(o(t)). h(c(t)))$$

$$= \frac{\partial \text{diag}(o(t))}{\partial w_{rec_3}} h(c(t)) + \text{diag}(o(t)) \cdot \frac{\partial h(c(t))}{\partial w_{rec_3}}$$

$\underbrace{\partial \text{diag}(o(t))}_{\Theta \leftarrow TR} \quad \underbrace{h(c(t)) \cdot \frac{\partial h(c(t))}{\partial w_{rec_3}}}_{(1 - c(t)) \circ (1 + c(t))}$

To find $\frac{\partial c(t)}{\partial w_{rec_3}}$ we re-write equation for $c(t)$ (eq. ③) as follows:

$$c(0) = 0 ; c(1) = c(0) + i(1) \circ \hat{y}_g(1) = i(1) \circ \hat{y}_g(1)$$

$$c(2) = c(1) + i(2) \circ \hat{y}_g(2)$$

\vdots

$$\boxed{c(t) = \sum_{k=1}^t i(k) \circ \hat{y}_g(k)}$$

P. 3

$$\Rightarrow c(t) = \sum_{k=1}^t \text{diag}(\hat{y}_g(k)) \cdot i(k)$$

$$\frac{\partial c(t)}{\partial w_{rec_3}} = \sum_{k=1}^t \left[\frac{\partial \text{diag}(\hat{y}_g(k))}{\partial w_{rec_3}} i(k) + \text{diag}(\hat{y}_g(k)) \cdot \frac{\partial i(k)}{\partial w_{rec_3}} \right]$$

$\underbrace{\Theta \leftarrow TR}_{\Theta \leftarrow TR}$

$$\Rightarrow \boxed{\frac{\partial c(t)}{\partial w_{rec_3}} = \sum_{k=1}^t \text{diag}(\hat{y}_g(k)) \cdot i(k) \circ (1 - i(k)) \cdot \hat{y}_g^T(k-1)} \quad \text{eq. (25)}$$

$$\Rightarrow \frac{\partial \hat{y}(t)}{\partial w_{rec_3}} = \text{diag}(a(t)) \cdot (1 - c(t)) \circ (1 + c(t)) \cdot \sum_{k=1}^t \text{diag}(\hat{y}_g(k)) i(k) \circ (1 - i(k)) \cdot \hat{y}_g^T(k-1)$$

$$\Rightarrow \frac{\partial \hat{y}(t)}{\partial w_{rec_3}} = \sum_{k=1}^t [a(t) \circ (1 - c(t)) \circ (1 + c(t)) \cdot \hat{y}_g(k) \circ i(k) \circ (1 - i(k))] \hat{y}_g^T(k-1)$$

$\cancel{\text{eq. (25)}} \quad \cancel{\text{eq. (24)}}$

But this is expensive to implement, we use the following trick:

$$\frac{\partial \hat{y}(t)}{\partial w_{rec_3}} = \sum_{k=1}^{t-1} [a(t) \circ b(k)] \hat{y}_g^T(k-1) + [a(t) \circ b(t)] \hat{y}_g^T(t-1)$$

Expensive Part.

Using Lemma ① we can re-write the expensive parts P. 6

$$\frac{\partial \hat{y}(t)}{\partial w_{rec_3}} = \text{diag}(a(t)) \sum_{k=1}^{t-1} b(k) \hat{y}_g^T(k-1) + [a(t) \circ b(t)] \hat{y}_g^T(t-1)$$

$$\frac{\partial \hat{y}(t)}{\partial w_{rec_3}} = \text{diag}(a(t)) \cdot \frac{\partial c(t-1)}{\partial w_{rec_3}} + [a(t) \circ b(t)] \hat{y}_g^T(t-1)$$

$$a(t) = o(t) \circ (1 - c(t)) \circ (1 + c(t))$$

$$b(t) = \hat{y}_g(t) \circ i(t) \circ (1 - i(t))$$

More ~~similar~~ Similar to eq. (25) and more understandable!!

$$\frac{\partial \hat{y}(t)}{\partial w_{rec_3}} = [\text{diag}(a(t)) \cdot \left[\frac{\partial c(t-1)}{\partial w_{rec_3}} + b(t) \hat{y}_g^T(t-1) \right]]$$

MANUAL GRADIENT CALCULATION

- Manual gradient calculation & implementation is prone to bugs, make sure to perform the “**gradient check**”

$$\frac{\partial \hat{y}(t)}{\partial w_{rec_3}} = \frac{\partial}{\partial w_{rec_3}} (\text{diag}(o(t)). h(c(t)))$$

$$= \frac{\partial \text{diag}(o(t))}{\partial w_{rec_3}} h(c(t)) + \text{diag}(o(t)) \cdot \frac{\partial h(c(t))}{\partial w_{rec_3}}$$

$\underbrace{\quad}_{\textcircled{1} \leftarrow \text{TR}}$

$$\underbrace{h'(c(t)) \cdot \frac{\partial c(t)}{\partial w_{rec_3}}}_{(1 - c(t)) \circ (1 + c(t))}$$

To find $\frac{\partial c(t)}{\partial w_{rec_3}}$ we re-write equation for $c(t)$ (eq. ③) as follows:

$$c(0) = 0 ; c(1) = c(0) + i(1) \circ \hat{y}_g(1) = i(1) \circ \hat{y}_g(1)$$

$$c(2) = c(1) + i(2) \circ \hat{y}_g(2)$$

\vdots

$$\left\{ c(t) = \sum_{k=1}^t i(k) \circ \hat{y}_g(k) \right\}$$

$$\Rightarrow c(t) = \sum_{k=1}^t \text{diag}(\hat{y}_g(k)) \cdot i(k)$$

P. 3

$$\frac{\partial c(t)}{\partial w_{rec_3}} = \sum_{k=1}^t \left[\frac{\partial \text{diag}(\hat{y}_g(k))}{\partial w_{rec_3}} i(k) + \text{diag}(\hat{y}_g(k)) \cdot \frac{\partial i(k)}{\partial w_{rec_3}} \right]$$

$\underbrace{\quad}_{\textcircled{2} \leftarrow \text{TR}}$

$$\Rightarrow \frac{\partial c(t)}{\partial w_{rec_3}} = \sum_{k=1}^t \text{diag}(\hat{y}_g(k)) \cdot i(k) \circ (1 - i(k)) \cdot \hat{y}_g^{T}(k-1)$$

eq. (25)

$$\Rightarrow \frac{\partial \hat{y}(t)}{\partial w_{rec_3}} = \text{diag}(a(t)) \cdot (1 - c(t)) \circ (1 + c(t)) \cdot \sum_{k=1}^t \text{diag}(\hat{y}_g(k)) i(k) \circ (1 - i(k)) \cdot \hat{y}_g^{T}(k-1)$$

$$\qquad \qquad \qquad \hat{y}_g(k) \circ i(k) \circ (1 - i(k))$$

$$\Rightarrow \frac{\partial \hat{y}(t)}{\partial w_{rec_3}} = \sum_{k=1}^t \left[\frac{a(t) \circ (1 - c(t)) \circ (1 + c(t))}{a(t)} \cdot \hat{y}_g(k) \circ i(k) \circ (1 - i(k)) \right] \hat{y}_g^{T}(k-1)$$

eq. (25)
eq. (24)

But this is expensive to implement, we use the following trick:

$$\frac{\partial \hat{y}(t)}{\partial w_{rec_3}} = \sum_{k=1}^{t-1} [a(t) \circ b(k)] \hat{y}_g^{T}(k-1) + [a(t) \circ b(t)] \hat{y}_g^{T}(t-1)$$

$\underbrace{\quad}_{\text{Expensive Part.}}$

Using lemma ① we can re-write the expensive parts P. 6

$$\frac{\partial \hat{y}(t)}{\partial w_{rec_3}} = \text{diag}(a(t)) \sum_{k=1}^{t-1} b(k) \hat{y}_g^{T}(k-1) + [a(t) \circ b(t)] \hat{y}_g^{T}(t-1)$$

$\underbrace{\quad}_{\frac{\partial c(t-1)}{\partial w_{rec_3}}}$

$$\Rightarrow \frac{\partial \hat{y}(t)}{\partial w_{rec_3}} = \text{diag}(a(t)) \cdot \frac{\partial c(t-1)}{\partial w_{rec_3}} + [a(t) \circ b(t)] \hat{y}_g^{T}(t-1)$$

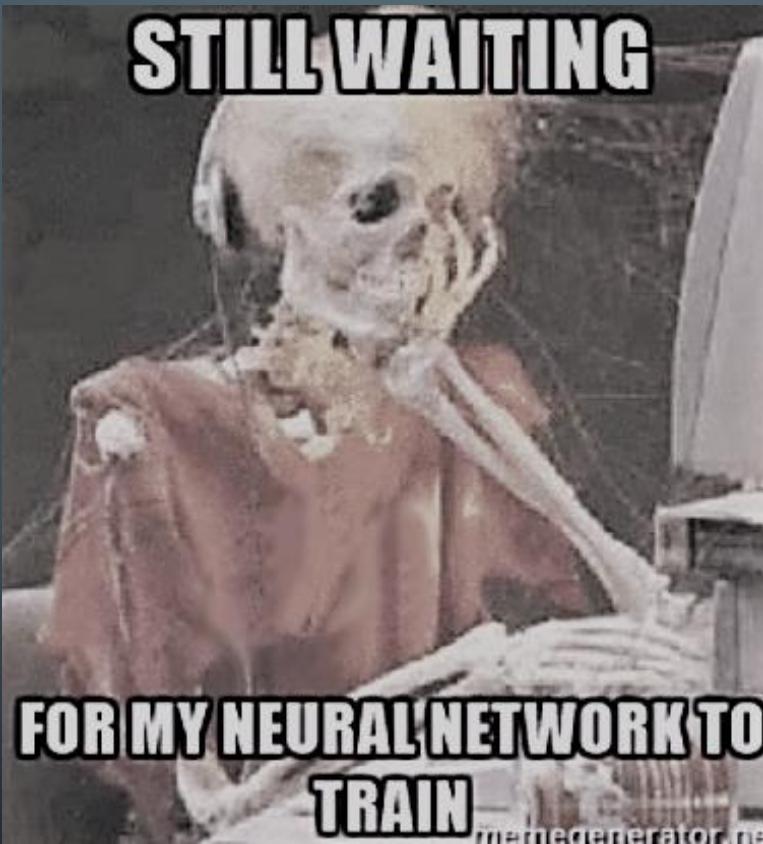
$$a(t) = o(t) \circ (1 - c(t)) \circ (1 + c(t))$$

$$b(t) = \hat{y}_g(t) \circ i(t) \circ (1 - i(t))$$

More similar to eq. (25) and more understandable!!

$$\Rightarrow \frac{\partial \hat{y}(t)}{\partial w_{rec_3}} = [\text{diag}(a(t))] \cdot \left[\frac{\partial c(t-1)}{\partial w_{rec_3}} + b(t) \hat{y}_g^{T}(t-1) \right]$$

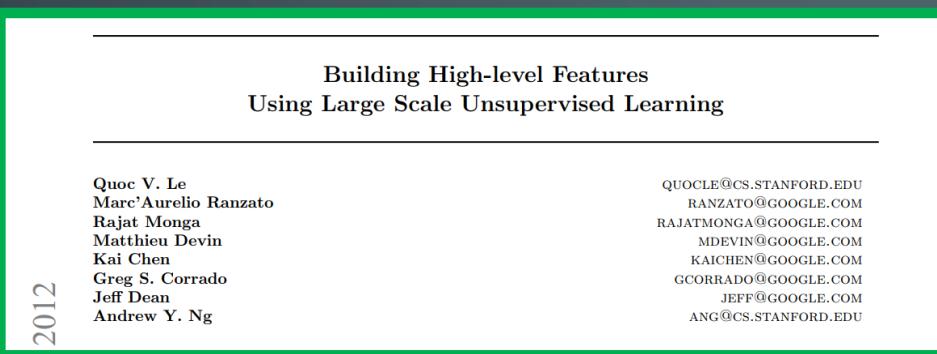
CPU_s and GPU_s



Picture from <https://www.analyticsvidhya.com/blog/2017/05/gpus-necessary-for-deep-learning/>

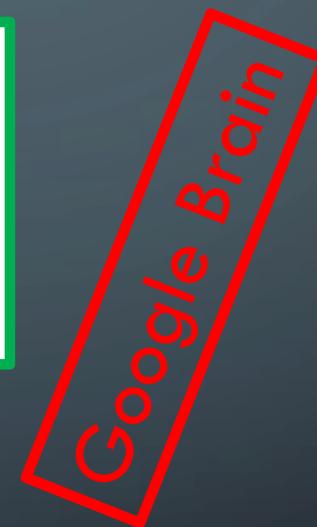
CPU_s and GPU_s

- [2012]



From <https://arxiv.org/abs/1112.6209>

2,000 CPUs (16,000 cores) – 600 kWatts - \$5,000,000



From <https://www.wired.com/2012/06/google-x-neural-network/>

CPU_s and GPU_s

- [2013]

3 GPUs (18,432 cores) – 4 kWatts - \$33,000

Stanford AI Lab (Andrew Ng)

DANIELA HERNANDEZ BUSINESS 06.17.13 6:30 AM

NOW YOU CAN BUILD GOOGLE'S ARTIFICIAL BRAIN ON THE CHEAP

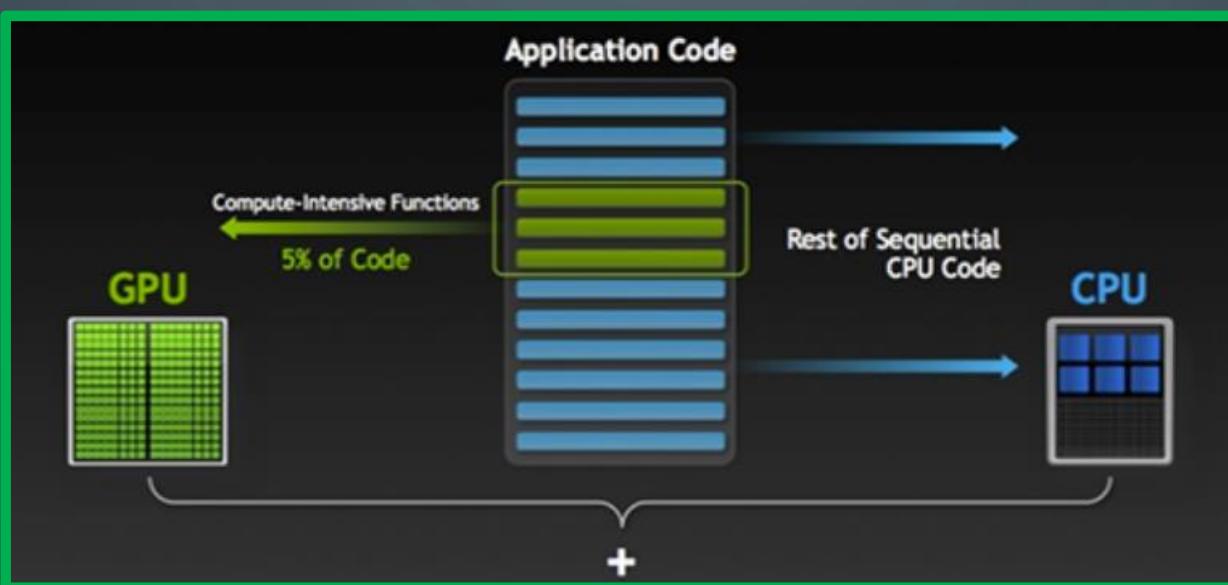


Andrew Ng. Photo: Ariel Zambelich/Wired

From https://www.wired.com/2013/06/andrew_ng/

CPU_s and GPU_s

- GPU acceleration



From <http://www.nvidia.com/object/what-is-gpu-computing.html>

CPUs and GPUs: (A SHORT FLAVOR OF) CUDA

- CUDA is C with a few extensions

- Use of function type qualifiers (`__global__`, `__device__`, `__host__`) to:
 - Determine if a function is executed on the host (CPU) or device (GPU)
 - Determine if a function is callable from the host or the device
- Use of variable type qualifiers (`__shared__`, `__device__`) to:
 - Determine the memory location of a variable
- Adding a new directive to:
 - Determine how a “**kernel**” is executed on the device from the host
- Using 4 built in variables (`gridDim`, `blockDim`, `blockIdx`, `threadIdx`) to:
 - Specify grid dimensions, block dimensions, block indices and thread indices

Called & executed by device

Called by host, executed by device

Called & executed by host

Variable in shared memory, has lifetime of block

A function that is executing portion of an application on the device (GPU)

Variable in global memory, has lifetime of the application

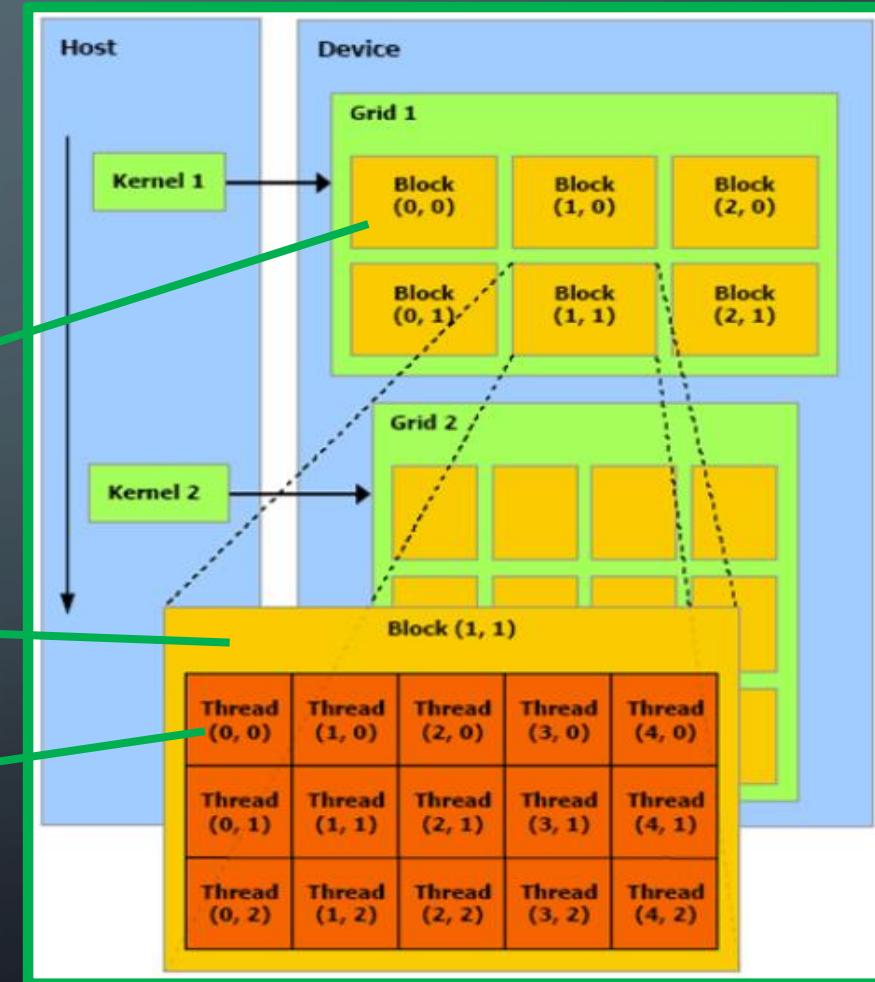
CPU_s and GPU_s: (A SHORT FLAVOR OF) CUDA

- A little more details

Blocks are independent, they must be able to execute in any order.

Threads in a block can synchronize execution

Kernels are executed by threads. Each thread has its own ID. Usually many threads execute the same kernel



CPUs and GPUs: (A SHORT FLAVOR OF) CUDA

- A little more details

Quite fast, R/W, Only accessed by 1 thread – This is thread space

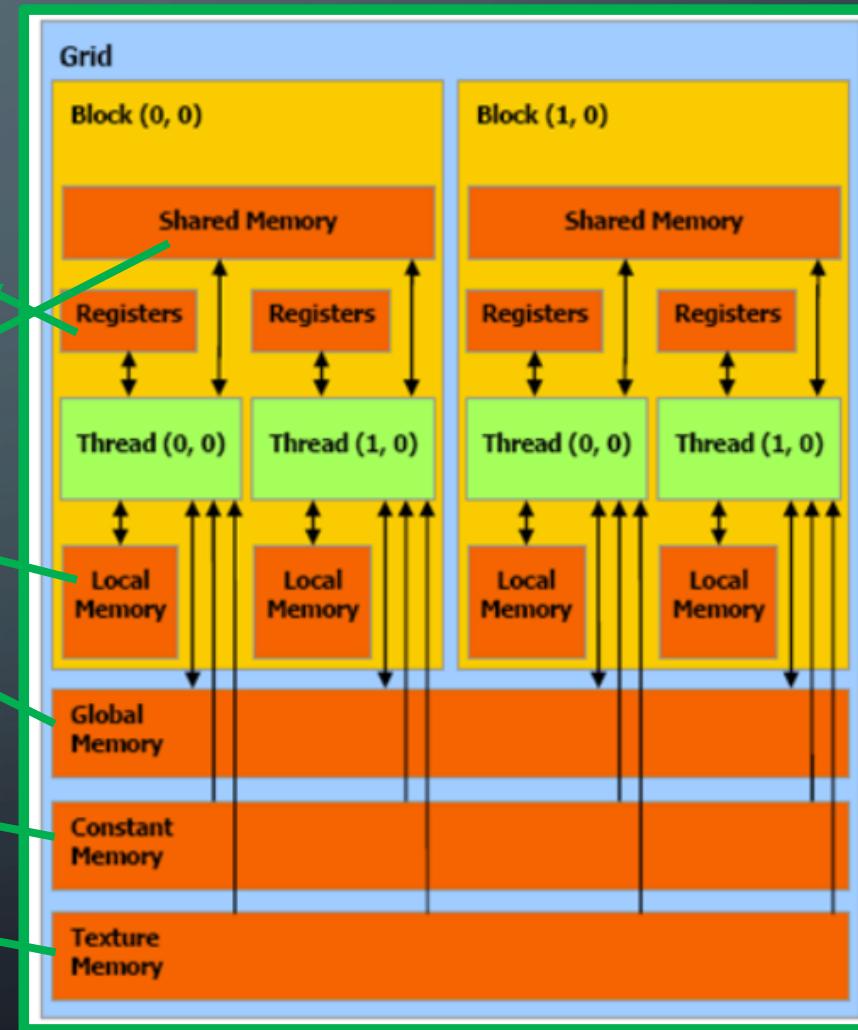
Fast, R/W, Accessed by all threads in a block (16 KB) – This is for thread collaboration

Quite fast, R/W, Only accessed by 1 thread

Not as fast as local & shared memory, R/W, Accessed by all threads and CPU (4 GB) – Used for IO for Grid

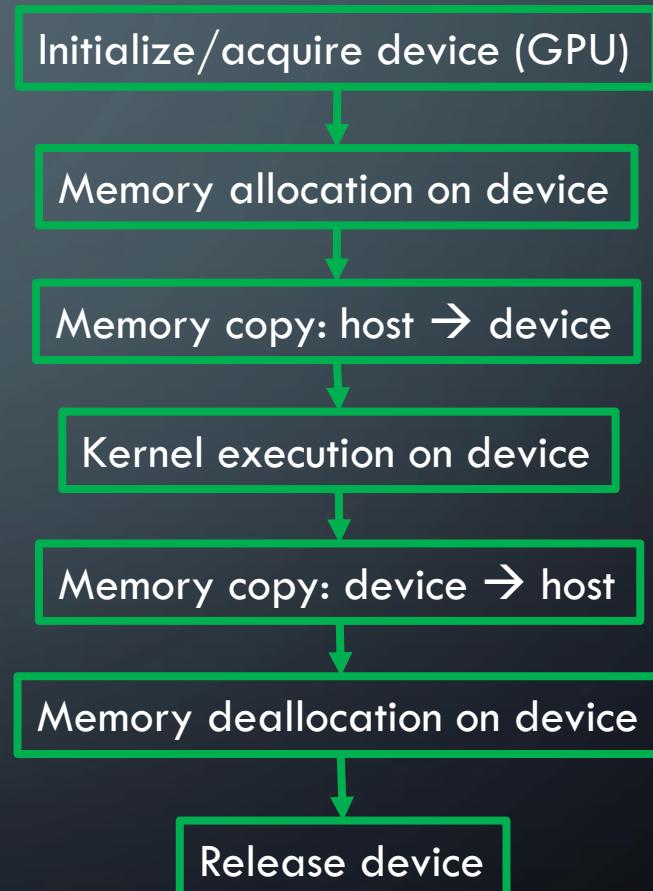
R, Accessed by all threads and CPU

R, Accessed by all threads and CPU



CPUs and GPUs: (A SHORT FLAVOR OF) CUDA

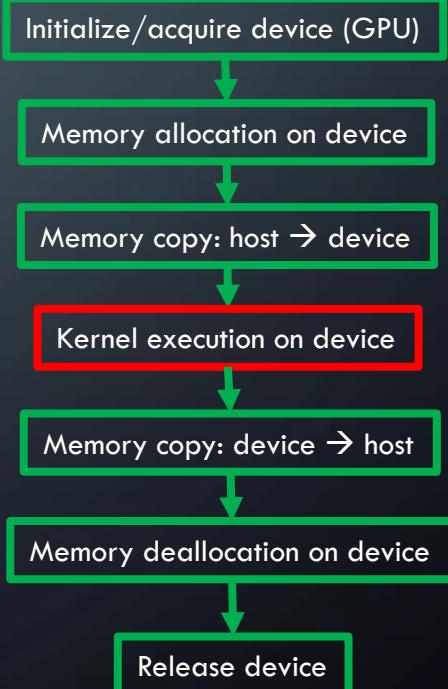
- Workflow [from http://geco.mines.edu/tesla/cuda_tutorial_mio/]



CPUs and GPUs: (A SHORT FLAVOR OF) CUDA

- Simple example: adding 2 arrays [from https://developer.download.nvidia.com/compute/DevZone/docs/html/C/doc/CUDA_C_Programming_Guide.pdf]

```
// Device code
__global__ void VecAdd(float* A, float* B, float* C, int N)
{
    int i = blockDim.x * blockIdx.x + threadIdx.x;
    if (i < N)
        C[i] = A[i] + B[i];
}
```



CPUs and GPUs: (A SHORT FLAVOR OF) CUDA

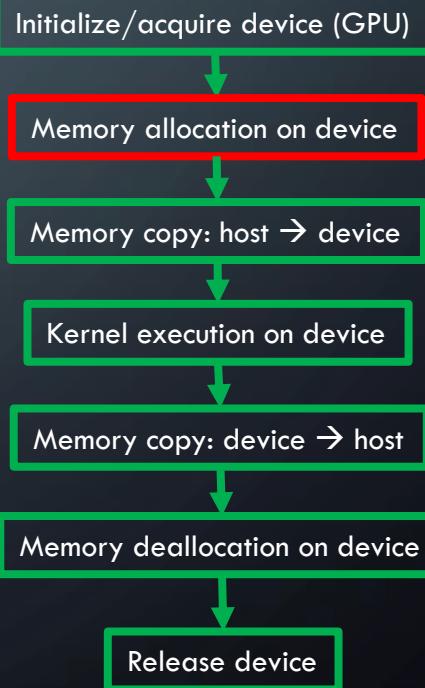
- Simple example: adding 2 arrays [from https://developer.download.nvidia.com/compute/DevZone/docs/html/C/doc/CUDA_C_Programming_Guide.pdf]

```
// Host code
int main()
{
    int N = ...;
    size_t size = N * sizeof(float);

    // Allocate input vectors h_A and h_B in host memory
    float* h_A = (float*)malloc(size);
    float* h_B = (float*)malloc(size);

    // Initialize input vectors
    ...

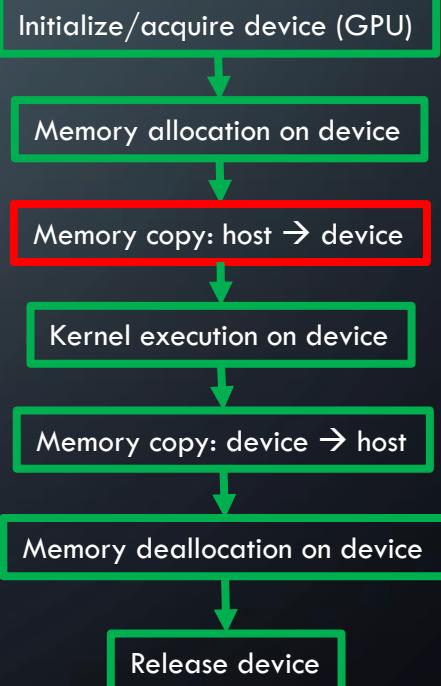
    // Allocate vectors in device memory
    float* d_A;
    cudaMalloc(&d_A, size);
    float* d_B;
    cudaMalloc(&d_B, size);
    float* d_C;
    cudaMalloc(&d_C, size);
```



CPUs and GPUs: (A SHORT FLAVOR OF) CUDA

- Simple example: adding 2 arrays [from https://developer.download.nvidia.com/compute/DevZone/docs/html/C/doc/CUDA_C_Programming_Guide.pdf]

```
// Copy vectors from host memory to device memory  
cudaMemcpy(d_A, h_A, size, cudaMemcpyHostToDevice);  
cudaMemcpy(d_B, h_B, size, cudaMemcpyHostToDevice);  
  
// Invoke kernel
```

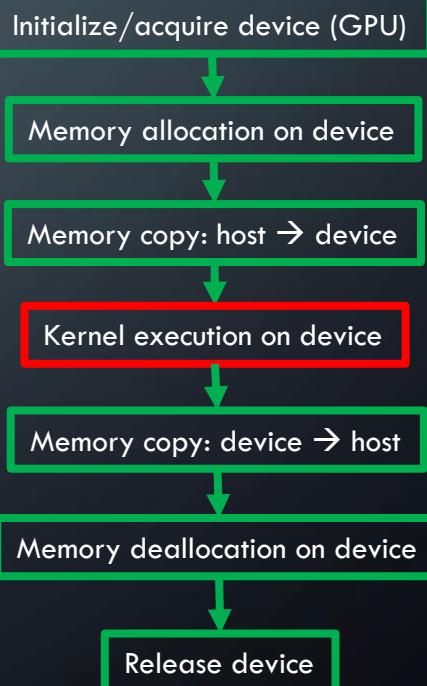


CPUs and GPUs: (A SHORT FLAVOR OF) CUDA

- Simple example: adding 2 arrays [from https://developer.download.nvidia.com/compute/DevZone/docs/html/C/doc/CUDA_C_Programming_Guide.pdf]

```
int threadsPerBlock = 256;
int blocksPerGrid =
    (N + threadsPerBlock - 1) / threadsPerBlock;
VecAdd<<<blocksPerGrid, threadsPerBlock>>>(d_A, d_B, d_C, N);

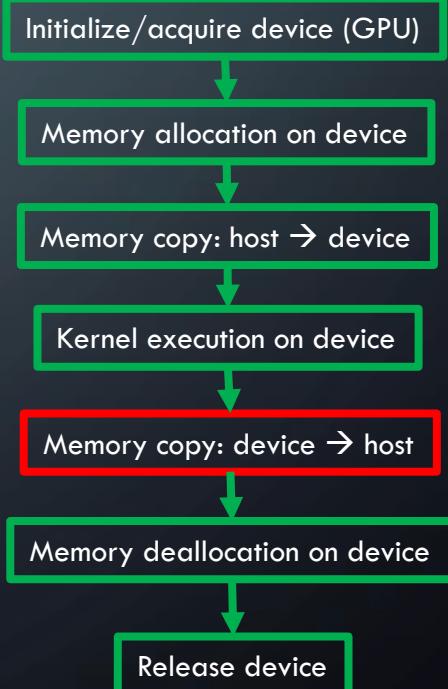
// Device code
__global__ void VecAdd(float* A, float* B, float* C, int N)
{
    int i = blockDim.x * blockIdx.x + threadIdx.x;
    if (i < N)
        C[i] = A[i] + B[i];
}
```



CPUs and GPUs: (A SHORT FLAVOR OF) CUDA

- Simple example: adding 2 arrays [from https://developer.download.nvidia.com/compute/DevZone/docs/html/C/doc/CUDA_C_Programming_Guide.pdf]

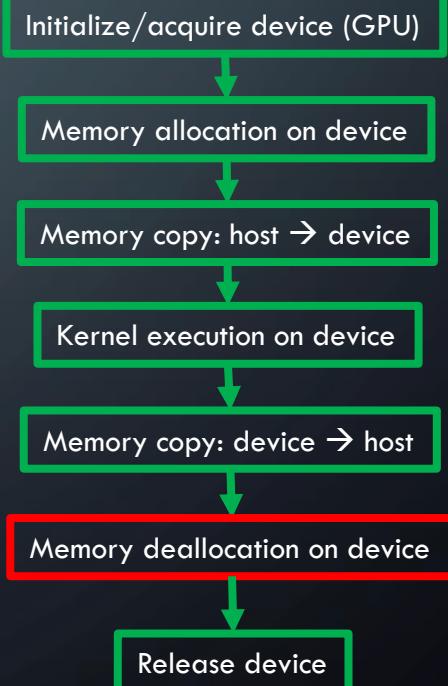
```
// Copy result from device memory to host memory  
// h_C contains the result in host memory  
cudaMemcpy(h_C, d_C, size, cudaMemcpyDeviceToHost);
```



CPUs and GPUs: (A SHORT FLAVOR OF) CUDA

- Simple example: adding 2 arrays [from https://developer.download.nvidia.com/compute/DevZone/docs/html/C/doc/CUDA_C_Programming_Guide.pdf]

```
// Free device memory  
cudaFree(d_A);  
cudaFree(d_B);  
cudaFree(d_C);  
  
// Free host memory  
...  
}
```



DL FRAMEWORKS: MICROSOFT COGNITIVE TOOLKIT (CNTK)

- From Microsoft
- Supported interfaces: C#, Python, C++ and Command Line
- High scalability
 - Scales across GPUs & machines
- Very fast for sequential models
 - E.g., RNNs, LSTMs
- No commercial support

DL FRAMEWORKS: TENSORFLOW

- From Google
- Supported interfaces: Python, C++ (and experimental support for Java API: not quite stable yet)
- Capability to run on multiple CPUs / GPUs.
- Computation graph compiles faster than Theano (RIP)
- There is no commercial support
- Creates static graphs
- Not closely similar to numpy

DL FRAMEWORKS: TORCH & PYTORCH

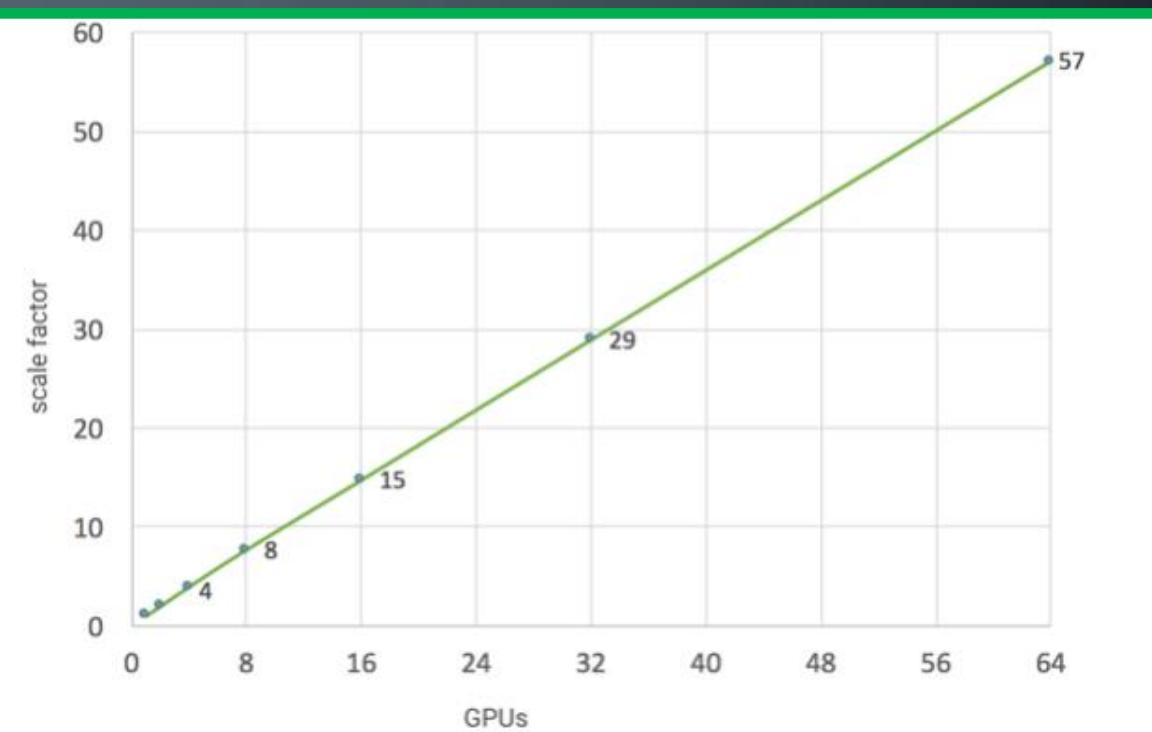
- Torch is Maintained by Facebook/Twitter/Google (DeepMind)
- Supported interfaces for Torch: C, C++, Lua
- PyTorch (open sourced in Jan. 2017 by Facebook) is not a Python binding into a C++ framework, it is built to be integrated in Python. Can be used naturally like numpy, scipy, ...
- PyTorch Tensors can be used either on CPU or GPU, a replacement for numpy to use GPUs
- PyTorch builds NNs dynamically [computation graph built at run-time]:
 - TensorFlow, CNTK, Caffe and Theano (RIP): Build NN & reuse it, if you want to change NN architecture, you should build another NN from scratch [static: computation graph is first “compiled”, and run after that]
 - PyTorch: Uses Reverse-mode auto-differentiation that allows changing NN’s behavior with quite low overhead = high flexibility for research projects
- It is easy to write your own layer types
- There is no commercial support

DL FRAMEWORKS: CAFFE

- From Berkeley Vision and Learning Center (BVLC)
- Supported interfaces: Python, MATLAB, C++, C, Command line
- Quite useful when using CNNs
- Rich set of pre-trained models (Model Zoo)
- Initial focus: Computer Vision
- Drawbacks: Documentation, many dependencies, flexibility (need to code in C++ and cuda for significant modifications to the architecture, e.g., new layers to be run on GPU)
- Appropriate for computer vision production code (robust and fast)
- For initial experiments and exploration use a high level API (e.g., Keras), after that use Caffe for production
- Not appropriate for Recurrent Neural Networks (RNNs)
- No commercial support

DL FRAMEWORKS: CAFFE2

- From Facebook, built on the top of caffe
- Supported interfaces: Python, C++
- High scaling properties
 - E.g., close to linear scaling with ResNet-50
 - Has made CNNs' distributed training easier
 - Better memory optimization than Caffe
 - Capability for mixed precision computations
 - float, float16, int8, ...
- No commercial support



ImageNet training using 64 NVIDIA Tesla P100 GPUs,
8 servers each one having 8 GPUs

[Figure from: <https://devblogs.nvidia.com/parallelforall/caffe2-deep-learning-framework-facebook/>]

DL FRAMEWORKS: THEANO (RIP)

- One of the first DL libraries, from Yoshua Bengio's lab (MILA)
- Supported interfaces: Python
- Not as scalable as other DL frameworks, e.g., lacks multi-GPU support
- A lot of low level coding should be done when using Theano (there are high level wrappers on the top of Theano though, e.g., Keras and Lasagne)
- Compile time of computation graph is too long sometimes
- On Sept. 28, 2017 MILA announced that it will stop developing Theano (RIP Theano)
...

DL FRAMEWORKS: MXNET

- Supported interfaces: Python, C++, R, Julia
- Scalable, can run experiments on multiple GPUs and machines
- Amazon's "DL framework of choice"

DL FRAMEWORKS: DEEPMLEARNING4J (DL4J)

- Developed by Skymind (a San Francisco-based software firm)
- Supported interfaces: Java & Scala, compatible with JVM
- Can be implemented on the top of Big Data tools, e.g., Apache Hadoop and Apache Spark
- Good documentation

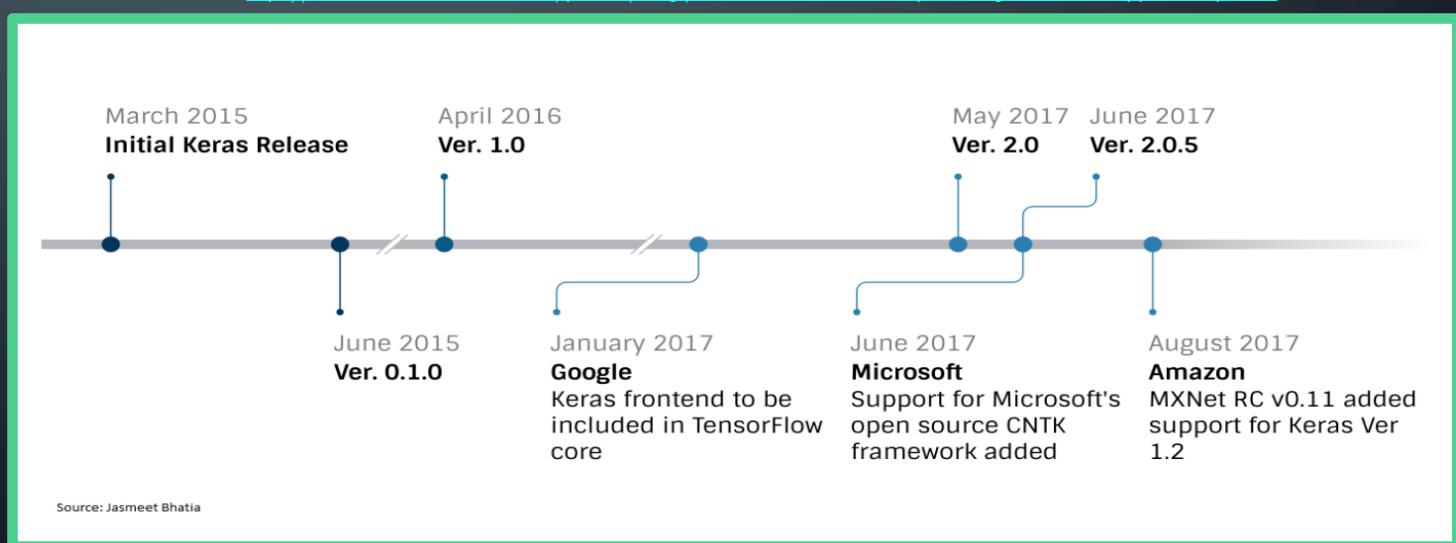
DL FRAMEWORKS: BIGDL

- From Intel
- Supported interfaces: Python and Scala
- Distributed DL library for Spark: Can run directly on the top of Apache Hadoop and Spark clusters.
- High scalability
- You can load pretrained Torch / Caffe models into Spark

DL FRAMEWORKS: HIGH LEVEL NN APIs

- An easier way to build your DL models:
 - Keras
 - Supported interface: Python
 - You can build a DL model in a few lines of code.
 - Can use Theano (RIP), TensorFlow, Microsoft Cognitive Toolkit (CNTK), MXNet or DL4j as backend
 - Good documentation

Picture from <https://www.datasciencecentral.com/profiles/blogs/search-for-the-fastest-deep-learning-framework-supported-by-keras>



DL FRAMEWORKS: HIGH LEVEL NN APIs

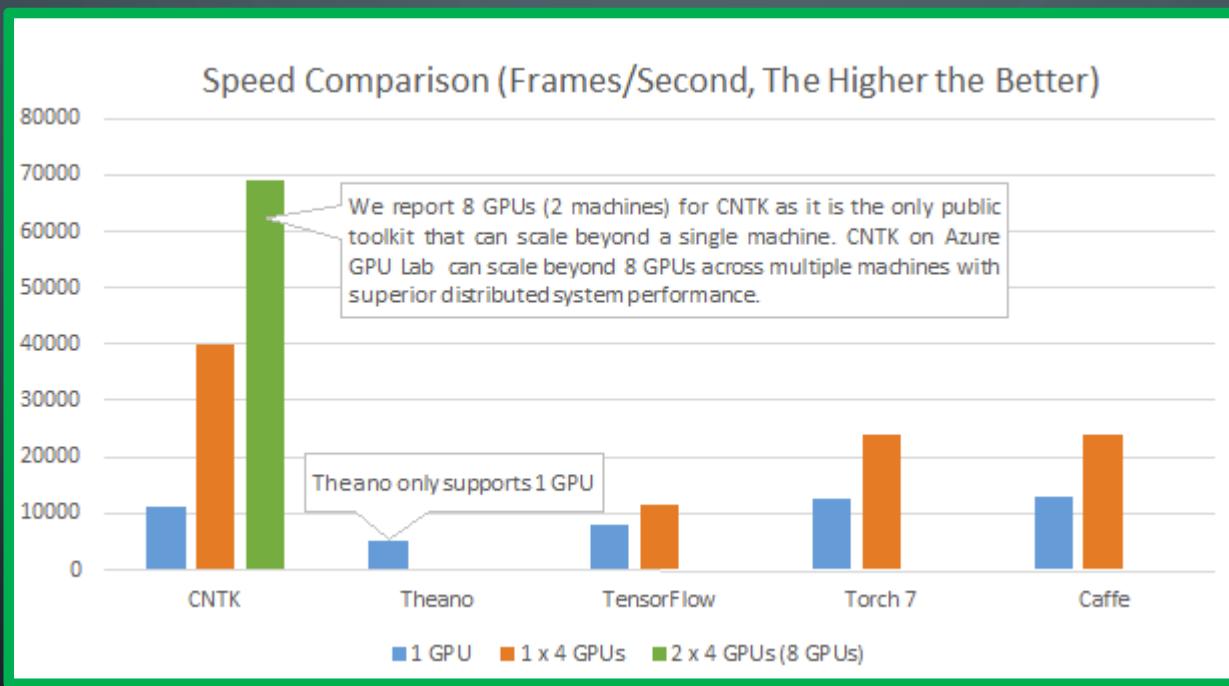
- An easier way to build your DL models:
 - Lasagne
 - Supported interface: Python
 - Can only use Theano (RIP) as backend
 - Not as good documentation as Keras
 - Note: although these APIs make it easy to build DL models, they might not be as flexible as using the backend DL libraries directly.

DL FRAMEWORKS: A FEW MORE

- Chainer
 - From a Tokyo start up named Preferred Networks
 - Was the only framework for dynamic computation graphs before PyTorch
- DSSTNE
 - From Amazon, written mainly in C++
 - Amazon decided to support MXNet on AWS
- DyNet
 - From CMU, supports dynamic computation graphs
 - Community is not as large as other frameworks
- Gluon
 - From Microsoft & Amazon: High level API on the top of MXNet [Announced Oct. 2017]
- Paddle
 - DL framework from Baidu

DL FRAMEWORKS: BENCHMARKING EFFORTS

- CNTK vs TensorFlow vs Theano vs Torch vs Caffe



Picture from MSR blog, Dec. 7, 2015

DL FRAMEWORKS: BENCHMARKING EFFORTS

- From <http://dlbench.comp.hkbu.edu.hk/> and <https://github.com/hclhkbu/dlbench>

alexnet on K80					
Batch Size	Caffe	CNTK	MXNET	TensorFlow	Torch
86	13.625ms(339.932s)	11.215ms(265.182s)	11.672ms(274.673s)	37.750ms(882.664s)	14.266ms(358.039s)
128	19.511ms(327.894s)	14.883ms(236.889s)	15.496ms(245.240s)	54.314ms(853.009s)	19.924ms(335.840s)
256	36.815ms(311.379s)	27.215ms(217.514s)	28.994ms(229.586s)	103.960ms(818.209s)	37.462ms(315.710s)
512	71.104ms(301.520s)	54.548ms(217.935s)	54.881ms(217.607s)	202.342ms(796.155s)	72.871ms(307.042s)
1024	140.007ms(297.183s)	103.069ms(206.114s)	105.975ms(210.728s)	398.114ms(783.322s)	143.613ms(302.626s)
2048	277.656ms(300.443s)	197.756ms(201.892s)	212.700ms(212.388s)	783.488ms(786.463s)	285.177ms(300.476s)
resnet on K80					
Batch Size	Caffe	CNTK	MXNET	TensorFlow	Torch
11	101.328ms(18984.640s)	77.144ms(14049.724s)	50.536ms(9191.549s)	108.827ms(19816.777s)	47.548ms(9141.084s)
16	109.304ms(14227.879s)	54.484ms(6831.130s)	59.298ms(7415.866s)	123.418ms(15453.923s)	58.778ms(7794.904s)
32	143.987ms(9561.758s)	81.470ms(5112.876s)	84.545ms(5287.561s)	181.404ms(11365.723s)	90.935ms(6021.554s)
64	225.325ms(7613.368s)	181.920ms(5710.661s)	147.274ms(4605.828s)	301.445ms(9453.074s)	164.761ms(5455.466s)
128	383.212ms(6553.839s)	288.178ms(2496.103s)	266.322ms(4164.944s)	563.190ms(8830.668s)	307.323ms(5079.835s)

DL FRAMEWORKS: BENCHMARKING EFFORTS

- From <http://dlbench.comp.hkbu.edu.hk/> and <https://github.com/hclhkbu/dlbench>

Item in cell: batchTime(totalTime)					
fcn5 on K80					
	Tesla K80, CUDA: 8.0 CUDNN: v5.1 CUDA_DRIVER: 367.48 Network: fcn5				
Batch Size	Caffe	CNTK	MXNET	TensorFlow	Torch
342	25.271ms(199.234s)	23.838ms(253.527s)	30.711ms(218.684s)	28.622ms(213.854s)	24.435ms(186.543s)
512	31.956ms(172.195s)	30.047ms(227.684s)	38.298ms(182.741s)	37.560ms(190.068s)	30.554ms(157.515s)
1024	55.329ms(151.994s)	51.038ms(205.680s)	60.448ms(144.837s)	62.044ms(159.047s)	52.154ms(135.299s)
2048	98.740ms(140.120s)	92.788ms(197.777s)	104.051ms(125.150s)	111.653ms(145.926s)	90.818ms(118.498s)
4096	184.107ms(132.149s)	175.332ms(190.883s)	182.601ms(110.287s)	211.366ms(138.959s)	168.122ms(110.418s)

Item in cell: batchTime(totalTime)					
Istm on K80					
	Tesla K80, CUDA: 8.0 CUDNN: v5.1 CUDA_DRIVER: 367.48 Network: Istm				
Batch Size	CNTK	MXNET	TensorFlow	Torch	
64	40.967ms(11922.492s)	87.171ms(1674.387s)	--	287.026ms(3412.095s)	
128	42.736ms(6230.025s)	151.189ms(1451.132s)	--	565.879ms(3388.249s)	
256	43.581ms(3187.478s)	288.142ms(1381.371s)	--	1130.606ms(3408.770s)	
512	49.712ms(1826.834s)	560.380ms(1344.809s)	--	2312.802ms(3547.659s)	
1024	63.695ms(1177.922s)	1101.696ms(1303.797s)	--	5073.561ms(4061.392s)	

DL FRAMEWORKS: BENCHMARKING EFFORTS

- Shaohuai Shi et al, IEEE CCBD 2016 [<https://arxiv.org/pdf/1608.07249v7.pdf>]

Time per mini-batch. Table from Shaohuai et al, 2016

		Desktop CPU (Threads used)				Server CPU (Threads used)						Single GPU		
		1	2	4	8	1	2	4	8	16	32	G980	G1080	K80
FCN-S	Caffe	1.324	0.790	0.578	15.444	1.355	0.997	0.745	0.573	0.608	1.130	0.041	0.030	0.071
	CNTK	1.227	0.660	0.435	-	1.340	0.909	0.634	0.488	0.441	1.000	0.045	0.033	0.074
	TF	7.062	4.789	2.648	1.938	9.571	6.569	3.399	1.710	0.946	0.630	0.060	0.048	0.109
	MXNet	4.621	2.607	2.162	1.831	5.824	3.356	2.395	2.040	1.945	2.670	-	0.106	0.216
	Torch	1.329	0.710	0.423	-	1.279	1.131	0.595	0.433	0.382	1.034	0.040	0.031	0.070
AlexNet-S	Caffe	1.606	0.999	0.719	-	1.533	1.045	0.797	0.850	0.903	1.124	0.034	0.021	0.073
	CNTK	3.761	1.974	1.276	-	3.852	2.600	1.567	1.347	1.168	1.579	0.045	0.032	0.091
	TF	6.525	2.936	1.749	1.535	5.741	4.216	2.202	1.160	0.701	0.962	0.059	0.042	0.130
	MXNet	2.977	2.340	2.250	2.163	3.518	3.203	2.926	2.828	2.827	2.887	0.020	0.014	0.042
	Torch	4.645	2.429	1.424	-	4.336	2.468	1.543	1.248	1.090	1.214	0.033	0.023	0.070
RenNet-50	Caffe	11.554	7.671	5.652	-	10.643	8.600	6.723	6.019	6.654	8.220	-	0.254	0.766
	CNTK	-	-	-	-	-	-	-	-	-	-	0.240	0.168	0.638
	TF	23.905	16.435	10.206	7.816	29.960	21.846	11.512	6.294	4.130	4.351	0.327	0.227	0.702
	MXNet	48.000	46.154	44.444	43.243	57.831	57.143	54.545	54.545	53.333	55.172	0.207	0.136	0.449
	Torch	13.178	7.500	4.736	4.948	12.807	8.391	5.471	4.164	3.683	4.422	0.208	0.144	0.523
FCN-R	Caffe	2.476	1.499	1.149	-	2.282	1.748	1.403	1.211	1.127	1.127	0.025	0.017	0.055
	CNTK	1.845	0.970	0.661	0.571	1.592	0.857	0.501	0.323	0.252	0.280	0.025	0.017	0.053
	TF	2.647	1.913	1.157	0.919	3.410	2.541	1.297	0.661	0.361	0.325	0.033	0.020	0.063
	MXNet	1.914	1.072	0.719	0.702	1.609	1.065	0.731	0.534	0.451	0.447	0.029	0.019	0.060
	Torch	1.670	0.926	0.565	0.611	1.379	0.915	0.662	0.440	0.402	0.366	0.025	0.016	0.051
AlexNet-R	Caffe	3.558	2.587	2.157	2.963	4.270	3.514	3.381	3.364	4.139	4.930	0.041	0.027	0.137
	CNTK	9.956	7.263	5.519	6.015	9.381	6.078	4.984	4.765	6.256	6.199	0.045	0.031	0.108
	TF	4.535	3.225	1.911	1.565	6.124	4.229	2.200	1.396	1.036	0.971	0.227	0.317	0.385
	MXNet	13.401	12.305	12.278	11.950	17.994	17.128	16.764	16.471	17.471	17.770	0.060	0.032	0.122
	Torch	5.352	3.866	3.162	3.259	6.554	5.288	4.365	3.940	4.157	4.165	0.069	0.043	0.141
RenNet-56	Caffe	6.741	5.451	4.989	6.691	7.513	6.119	6.232	6.689	7.313	9.302	-	0.116	0.378
	CNTK	-	-	-	-	-	-	-	-	-	-	0.206	0.138	0.562
	TF	-	-	-	-	-	-	-	-	-	-	0.225	0.152	0.523
	MXNet	34.409	31.255	30.069	31.388	44.878	43.775	42.299	42.965	43.854	44.367	0.105	0.074	0.270
	Torch	5.758	3.222	2.368	2.475	8.691	4.965	3.040	2.560	2.575	2.811	0.150	0.101	0.301
LSTM	Caffe	-	-	-	-	-	-	-	-	-	-	-	-	-
	CNTK	0.186	0.120	0.090	0.118	0.211	0.139	0.117	0.114	0.114	0.198	0.018	0.017	0.043
	TF	4.662	3.385	1.935	1.532	6.449	4.351	2.238	1.183	0.702	0.598	0.133	0.065	0.140
	MXNet	-	-	-	-	-	-	-	-	-	-	0.089	0.079	0.149
	Torch	6.921	3.831	2.682	3.127	7.471	4.641	3.580	3.260	5.148	5.851	0.399	0.324	0.560

DL FRAMEWORKS: BENCHMARKING EFFORTS

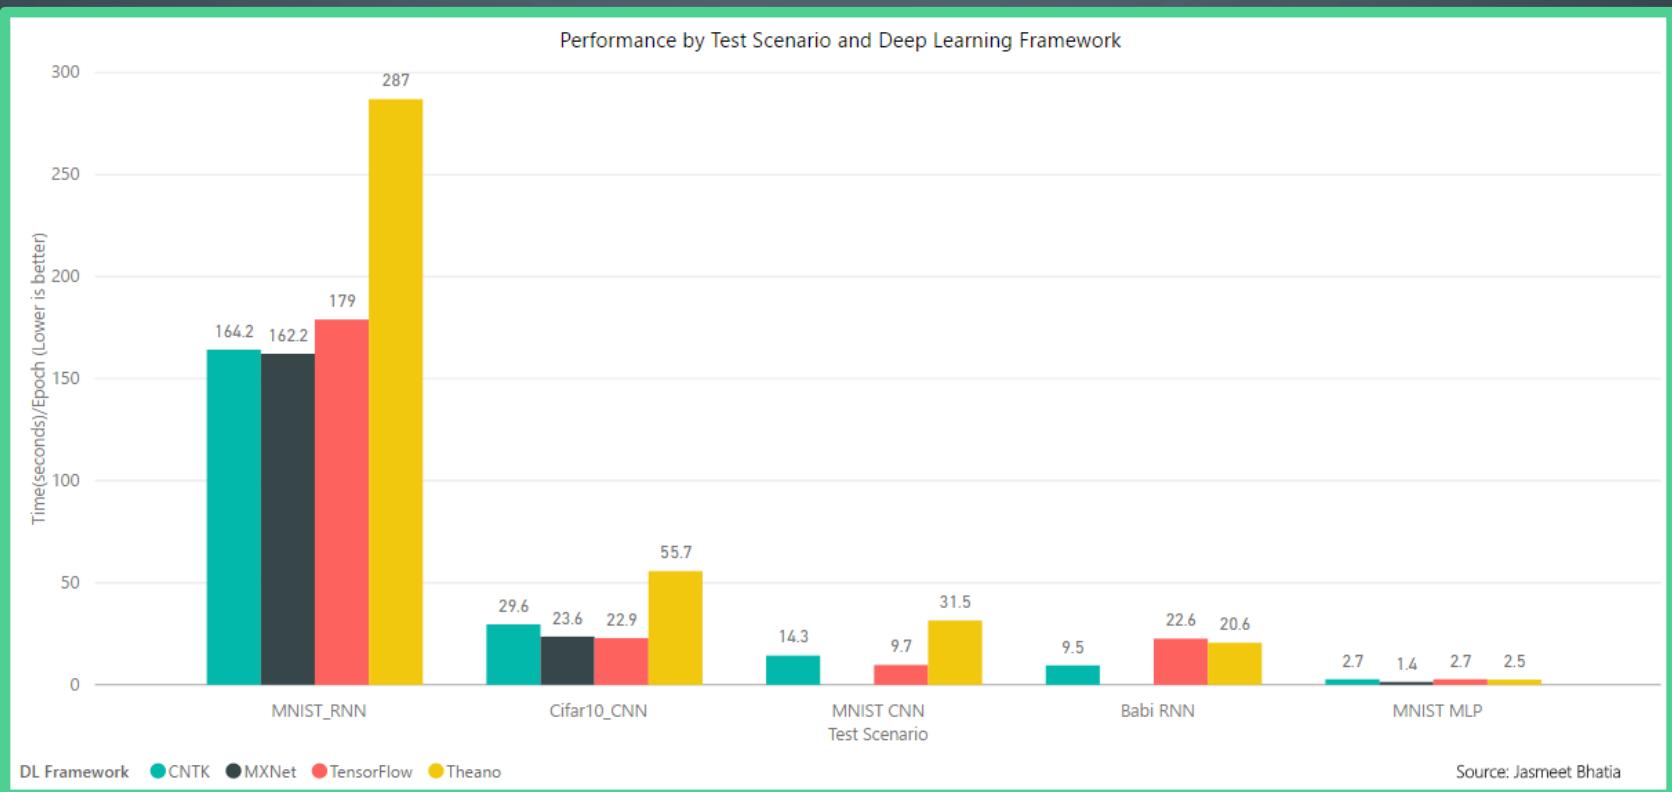
- TensorFlow vs MXNet on CIFAR-10, 8 GPUs [from <https://medium.com/@julsimon/keras-shoot-out-tensorflow-vs-mxnet-51ae2b30a9c0>]

```
2 Using MXNet backend.  
3 X_train shape: (50000, 3, 32, 32)  
4 50000 train samples  
5 10000 test samples  
6 [output removed]  
7 Epoch 2/100  
8 50000/50000 [=====] - 25s - loss: 6.7589 - acc: 0.4978 - val_loss: 7.0586  
9 Epoch 3/100  
10 50000/50000 [=====] - 25s - loss: 6.4458 - acc: 0.5799 - val_loss: 7.2990  
11 Epoch 4/100  
12 50000/50000 [=====] - 25s - loss: 6.2015 - acc: 0.6478 - val_loss: 7.3075  
13 Epoch 5/100  
14 50000/50000 [=====] - 25s - loss: 5.9922 - acc: 0.7095 - val_loss: 7.1960  
15 Epoch 6/100  
16 50000/50000 [=====] - 25s - loss: 5.8184 - acc: 0.7620 - val_loss: 6.9669  
17 Epoch 7/100  
18 50000/50000 [=====] - 25s - loss: 5.6454 - acc: 0.8120 - val_loss: 7.3894  
19 Epoch 8/100  
20 50000/50000 [=====] - 25s - loss: 5.5012 - acc: 0.8554 - val_loss: 7.2800  
21 Epoch 9/100  
22 50000/50000 [=====] - 25s - loss: 5.3884 - acc: 0.8847 - val_loss: 7.0814  
23 Epoch 10/100  
24 50000/50000 [=====] - 25s - loss: 5.2744 - acc: 0.9147 - val_loss: 7.2281
```

```
2 Using TensorFlow backend.  
3 X_train shape: (50000, 32, 32, 3)  
4 50000 train samples  
5 10000 test samples  
6 [output removed]  
7 2017-09-03 13:32:03.432572: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1030] Creating Tensor  
8 2017-09-03 13:32:03.432584: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1030] Creating Tensor  
9 2017-09-03 13:32:03.432600: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1030] Creating Tensor  
10 2017-09-03 13:32:03.432605: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1030] Creating Tensor  
11 2017-09-03 13:32:03.432611: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1030] Creating Tensor  
12 2017-09-03 13:32:03.432616: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1030] Creating Tensor  
13 2017-09-03 13:32:03.432620: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1030] Creating Tensor  
14 2017-09-03 13:32:03.432629: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1030] Creating Tensor  
15 [output removed]  
16 Epoch 2/100  
17 50000/50000 [=====] - 61s - loss: 6.8205 - acc: 0.4575 - val_loss: 6.8522  
18 Epoch 3/100  
19 50000/50000 [=====] - 61s - loss: 6.6435 - acc: 0.5241 - val_loss: 6.7707  
20 Epoch 4/100  
21 50000/50000 [=====] - 61s - loss: 6.5110 - acc: 0.5736 - val_loss: 6.7631  
22 Epoch 5/100  
23 50000/50000 [=====] - 61s - loss: 6.3973 - acc: 0.6166 - val_loss: 6.6761  
24 Epoch 6/100  
25 50000/50000 [=====] - 61s - loss: 6.2915 - acc: 0.6574 - val_loss: 6.6398  
26 Epoch 7/100  
27 50000/50000 [=====] - 61s - loss: 6.1942 - acc: 0.6957 - val_loss: 6.6293  
28 Epoch 8/100  
29 50000/50000 [=====] - 61s - loss: 6.1005 - acc: 0.7357 - val_loss: 6.6153  
30 Epoch 9/100  
31 50000/50000 [=====] - 61s - loss: 6.0076 - acc: 0.7704 - val_loss: 6.6278  
32 Epoch 10/100  
33 50000/50000 [=====] - 61s - loss: 5.9199 - acc: 0.8050 - val_loss: 6.7057
```

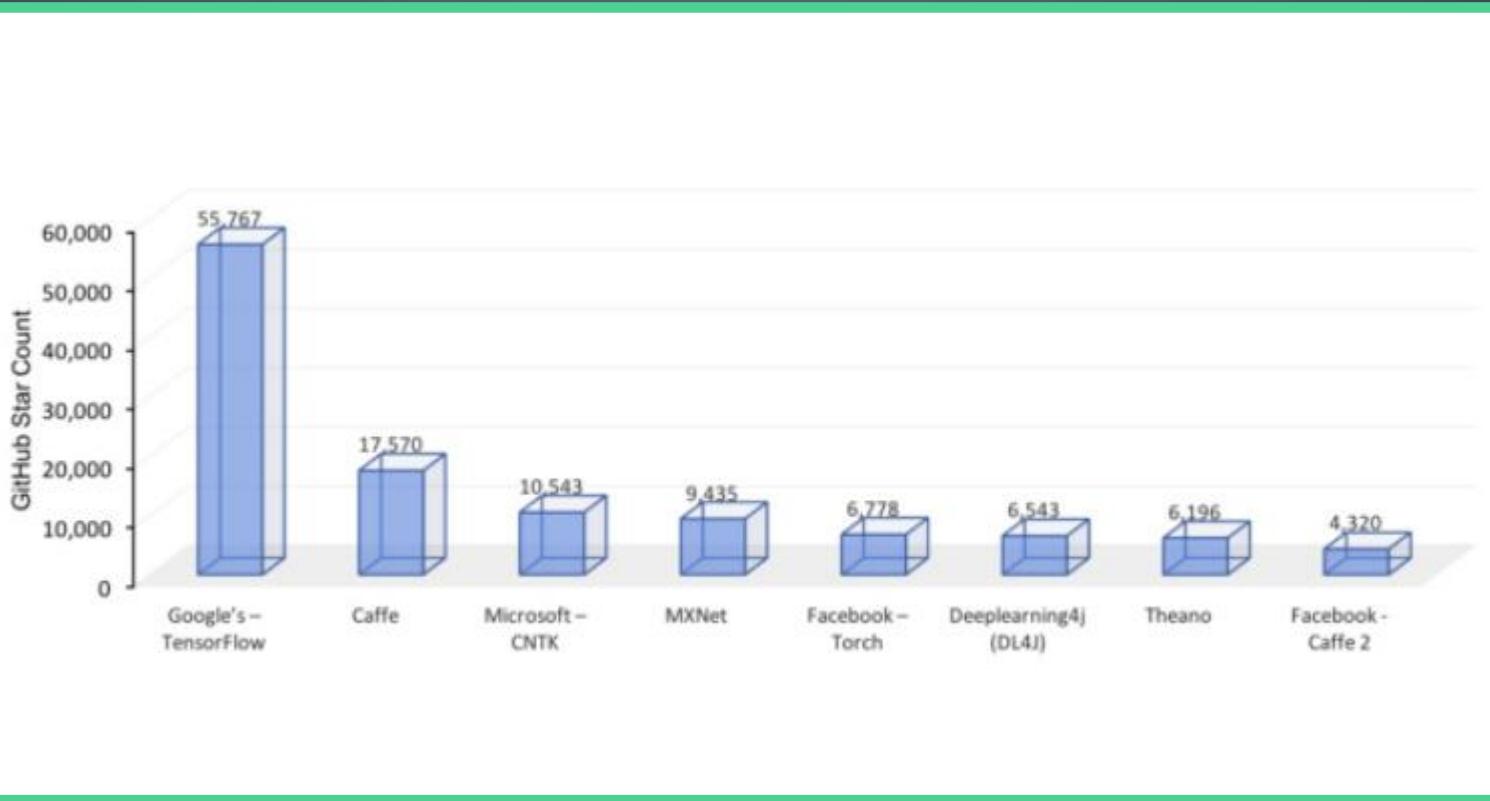
DL FRAMEWORKS: BENCHMARKING EFFORTS

- Benchmarking using Keras [graph from <https://www.datasciencecentral.com/profiles/blogs/search-for-the-fastest-deep-learning-framework-supported-by-keras>]



DL FRAMEWORKS: ON GITHUB

- GitHub star count! [from <https://www.cio.com/article/3193689/artificial-intelligence/which-deep-learning-network-is-best-for-you.html>]



DL FRAMEWORKS: PYTORCH VS TENSORFLOW

- Community of PyTorch is not as large as TensorFlow
- PyTorch does not have a visualization tool as powerful as Tensorboard in TensorFlow
- PyTorch is better for rapid prototyping for research. TensorFlow is better for large scale deployment
- PyTorch is easier to learn for beginners
- TensorFlow builds computation graphs “statically” but PyTorch does it “dynamically”

```
for _ in range(N):
    y = torch.matmul(W, y) + bias
```

DL FRAMEWORKS: PYTORCH VS TENSORFLOW

- PyTorch code is easier to debug than TensorFlow
- PyTorch is new & does not cover all functionalities yet (e.g., there is no fft in PyTorch yet). Over time with more contributions to PyTorch this gap will be closed ...
- TensorFlow is better for deployment
 - Using TensorFlow serialization the whole graph (including parameters and operations) can be saved, and then loaded for inference in other languages like C++ and Java. Quite useful when Python is not an option for deployment.
 - TensorFlow also works for mobile deployments (building mobile apps with TF <https://www.tensorflow.org/mobile/>), you don't need to code the DL architecture again for inference
- TensorFlow assumes you are using GPUs if any available. In PyTorch, everything should be explicitly moved to the device when using GPUs.
- PyTorch also have visualization tools like with similarities to TensorFlow's Tensorboard
 - Tensorboard_logger: https://github.com/TeamHG-Memex/tensorboard_logger
 - Crayon: <https://github.com/torrvision/crayon>

DL FRAMEWORKS

- From https://en.wikipedia.org/wiki/Comparison_of_deep_learning_software

Software	Creator	Software license ^[a]	Open source	Platform	Written in	Interface	OpenMP support	OpenCL support	CUDA support	Automatic differentiation ^[1]	Has pretrained models	Recurrent nets	Convolutional nets	RBM/DBNs	Parallel execution (multi node)
DeepLearning4j	Skymind engineering team; DeepLearning4j community; originally Adam Gibson	Apache 2.0	Yes	Linux, Mac OS X, Windows, Android (Cross-platform)	C++, Java	Java, Scala, Clojure, Python (Keras), Kotlin	Yes	On roadmap ^[8]	Yes ^{[9][10]}	Computational Graph	Yes ^[11]	Yes	Yes	Yes	Yes ^[12]
Dlib	Davis King	Boost Software License	Yes	Cross-Platform	C++	C++	Yes	No	Yes	Yes	Yes	No	Yes	Yes	Yes
Keras	François Chollet	MIT license	Yes	Linux, Mac OS X, Windows	Python	Python, R	Only if using Theano or MXNet as backend	Under development for the Theano backend (and on roadmap for the TensorFlow backend)	Yes	Yes	Yes ^[13]	Yes	Yes	Yes ^[14]	
MXNet	Distributed (Deep) Machine Learning Community	Apache 2.0	Yes	Linux, Mac OS X, Windows, AWS, Android, iOS, JavaScript ^[28]	Small C++ core library	C++, Python, Julia, Matlab, JavaScript, Go, R, Scala, Perl	Yes	On roadmap ^[29]	Yes	Yes ^[30]	Yes ^[31]	Yes	Yes	Yes	Yes ^[32]
Apache SINGA	Apache Incubator	Apache 2.0	Yes	Linux, Mac OS X, Windows	C++	Python, C++, Java	No	Yes	Yes	?	Yes	Yes	Yes	Yes	Yes
TensorFlow	Google Brain team	Apache 2.0	Yes	Linux, Mac OS X, Windows ^[33]	C++, Python	Python (Keras), C/C++, Java, Go, R ^[34]	No	On roadmap ^[35] but already with SYCL ^[36] support	Yes	Yes ^[37]	Yes ^[38]	Yes	Yes	Yes	Yes
Theano	Université de Montréal	BSD license	Yes	Cross-platform	Python	Python (Keras)	Yes	Under development ^[39]	Yes	Yes ^{[40][41]}	Through Lasagne's model zoo ^[42]	Yes	Yes	Yes	Yes ^[43]
Torch	Ronan Collobert, Koray Kavukcuoglu, Clement Farabet	BSD license	Yes	Linux, Mac OS X, Windows ^[44] , Android, iOS	C, Lua	Lua, LuaJIT, ^[45] C, utility library for C++/OpenCL ^[47]	Yes	Third party implementations ^{[48][49]}	Yes ^{[50][51]}	Through Twitter's Autograd ^[52]	Yes ^[53]	Yes	Yes	Yes	Yes ^[54]
Wolfram Mathematica	Wolfram Research	Proprietary	No	Windows, Mac OS X, Linux, Cloud computing	C++	Wolfram Language	No	No	Yes	Yes	Yes ^[55]	Yes	Yes	Yes	Yes
Microsoft Cognitive Toolkit	Microsoft Research	MIT license ^[16]	Yes	Windows, Linux ^[15] (OSX via Docker on roadmap)	C++	Python (Keras), C++, Command line ^[17] , BrainScript ^[18] (.NET on roadmap ^[19])	Yes ^[20]	No	Yes	Yes	Yes ^[21]	Yes ^[22]	Yes ^[22]	No ^[23]	Yes ^[24]
Caffe	Berkeley Vision and Learning Center	BSD license	Yes	Linux, Mac OS X, Windows ^[2]	C++	Python, MATLAB	Yes	Under development ^[3]	Yes	Yes	Yes ^[4]	Yes	Yes	No	?
Caffe2	Facebook	Apache 2.0	Yes	Linux, Mac OS X, Windows ^[5]	C++, Python	Python, MATLAB	Yes	Under development ^[6]	Yes	Yes	Yes ^[7]	Yes	Yes	No	Yes
MatConvNet	Andrea Vedaldi, Karen Lenc	BSD license	Yes	Windows, Linux ^[15] (OSX via Docker on roadmap)	C++	MATLAB, C++,	No	No	Yes	Yes	Yes	Yes	Yes	No	Yes
Neural Designer	Artelnics	Proprietary	No	Linux, Mac OS X, Windows	C++	Graphical user interface	Yes	No	No	?	?	No	No	No	?
OpenNN	Artelnics	GNU LGPL	Yes	Cross-platform	C++	C++	Yes	No	No	?	?	No	No	No	?
Gensim															
Paddle															
Pytorch															

DL FRAMEWORKS

- Microsoft Cognitive Toolkit (CNTK): <https://www.microsoft.com/en-us/cognitive-toolkit/>
- TensorFlow: <https://www.tensorflow.org/>
- Torch: <http://torch.ch/>
- PyTorch: <http://pytorch.org/>
- Caffe: <http://caffe.berkeleyvision.org/>
- Caffe2: <https://caffe2.ai/>
- Theano (RIP): <http://deeplearning.net/software/theano/>
- MXNet: <http://mxnet.incubator.apache.org/>
- Deeplearning4j: <https://deeplearning4j.org/>
- BigDL: <https://software.intel.com/en-us/articles/bigdl-distributed-deep-learning-on-apache-spark>
- High level NN APIs:
 - Keras [CNTK, TensorFlow, MXNet, DL4J and Theano backend]: <https://keras.io/>
 - Lasagne: [Theano backend] <https://lasagne.readthedocs.io/en/latest/>
- Chainer: <https://chainer.org/>
- DSSTNE: <https://github.com/amzn/amazon-dsstne>
- DyNet: <https://github.com/clab/dynet>
- Gluon: <https://github.com/gluon-api/gluon-api/>
- Paddle: <https://github.com/PaddlePaddle/Paddle>

DL FRAMEWORKS

- Microsoft Cognitive Toolkit (CNTK): <https://www.microsoft.com/en-us/cognitive-toolkit/>
- TensorFlow: <https://www.tensorflow.org/>
- Torch: <http://torch.ch/>
- PyTorch: <http://pytorch.org/>
- Caffe: <http://caffe.berkeleyvision.org/>
- Caffe2: <https://caffe2.ai/>
- Theano (RIP): <http://deeplearning.net/software/theano/>
- MXNet: <http://mxnet.incubator.apache.org/>
- Deeplearning4j: <https://deeplearning4j.org/>
- BigDL: <https://software.intel.com/en-us/articles/bigdl-distributed-deep-learning-on-apache-spark>
- High level NN APIs:
 - Keras [CNTK, TensorFlow, MXNet, DL4J and Theano backend]: <https://keras.io/>
 - Lasagne: [Theano backend] <https://lasagne.readthedocs.io/en/latest/>
- Chainer: <https://chainer.org/>
- DSSTNE: <https://github.com/amzn/amazon-dsstne>
- DyNet: <https://github.com/clab/dynet>
- Gluon: <https://github.com/gluon-api/gluon-api/>
- Paddle: <https://github.com/PaddlePaddle/Paddle>



DL FRAMEWORKS

No! Not Another Deep Learning Framework

Linh Nguyen*
lvnguyen@umich.edu
University of Michigan

Peifeng Yu*
peifeng@umich.edu
University of Michigan

Mosharaf Chowdhury
mosharaf@umich.edu
University of Michigan

Picture from <http://www.mosharaf.com/wp-content/uploads/deepstack-hotos17.pdf>

DL FRAMEWORKS: WHO ARE YOU?!

- Are you industry?
 - Speed and scale
 - Stability

DL4J

Microsoft Cognitive Toolkit (CNTK)

Caffe

TensorFlow

BigDL

MXNet ?!, Caffe2 ?!, DSSTNE ?!

DL FRAMEWORKS: WHO ARE YOU?!

- Are you a research organization?
 - Flexibility
 - Easy debuggability

PyTorch & Torch

Theano (RIP)

MXNet

TensorFlow

Microsoft Cognitive Toolkit (CNTK)

DL FRAMEWORKS: WHO ARE YOU?!

- Are you a DL beginner?
 - Do gradient calculation and backprop manually on paper once to fully understand it
 - Then start with a high level API to train your first DL model

Keras

Lasagne

DL FRAMEWORKS: WHO ARE YOU?!

- Are you a DL practitioner who wants to implement a model ASAP?
 - Use a high level API

Keras

Lasagne

DL FRAMEWORKS: WHO ARE YOU?!

- Are you a university prof planning to use a DL framework for your class?
 - Use an easy to learn framework with fast ramp-up time

PyTorch

MXNet

TensorFlow

Microsoft Cognitive Toolkit (CNTK)

DL FRAMEWORKS: WHO ARE YOU?!

- Are you an organization or company that needs commercial support?

DL4J

DL FRAMEWORKS: WHO ARE YOU?!

- Are you doing computer vision?

Caffe

Caffe2

MXNet

Torch

Microsoft Cognitive Toolkit (CNTK)

DL FRAMEWORKS: WHO ARE YOU?!

- Are you using RNNs (LSTMs, GRUs, ...) and variable length sequences?

Microsoft Cognitive Toolkit (CNTK)

PyTorch



THANK YOU!